

# Chapter 4

## Concurrency Control II

---

Multi-Granularity Locking  
Multi-Version Locking  
Predicate Locks  
Semantic Synchronization  
Escrow



### Multi-Granularity Locking (1)

---


#### ▪ Lock Granulate

- Determines parallelism/overhead
- Fine granulate reduces lock conflicts; however, many locks need to be acquired and managed
- Hierarchical processing allows flexibility w.r.t. granulate ('multi-granularity locking')
- e. g. Synchronization of long TA at table level
- or short TA at row level
- Commercial DBS mostly support at least 2 levels, e.g.
  - Page – Segment
  - Record type (Table) – Record (Row)



## Multi-Granularity Locking (2)

- Not restricted to hierarchies, can be extended to partially ordered object sets




- More complex than simple concurrency control mechanisms (more lock modes, conversion, deadlock-handling, ...)



## Multi-Granularity Locking (3)

- Lock hierarchy example**



Overhead of locking ...

- 1 record : 3 + 1
- k records : 3 + k
- 1 record type : 2 + 1




## Multi-Granularity Locking (4)

### Intention Locks

- R- and X-locks also lock all successor nodes implicitly
- All predecessor nodes are to be locked, too, in order to avoid incompatibilities
- Exploitation of so called 'intention locks'
- General intention lock: I-lock

	I	R	X
I	+	-	-
R	-	+	-
X	-	-	-




## Multi-Granularity Locking (5)

### Intention Locks (contd.)

- General intention lock – example
- Incompatibility of I- and R-locks
  - too restrictive!
- Solution (?): 2 intention lock modes: IR und IX


	IR	IX	R	X
IR	+	+	+	-
IX	+	+	-	-
R	+	-	+	-
X	-	-	-	-



## Multi-Granularity Locking (6)

### Intention Locks (contd.)

- IR and IX - Example
- IR-lock (intention read), if only read access on lower objects, otherwise IX-lock
- Further refinement:  $RIX = R + IX$ 
  - For the case that all records of a record type are supposed to be read and only some of them to be written
    - X-lock on record type would be too restrictive
    - IX-lock on record type would require to lock each respective record
  - Locks object in R-mode and requires ...
  - ... X-locks at lower hierarchy level only for objects which are to be updated




## Multi-Granularity Locking (7)

### Intention Locks (contd.)

- Complete protocol
  - RIX allows read access on the node and its successors; further, it encompasses the right to acquire IX-, U- and X-locks on successors
  - U: read with write intention; conversion  $U \rightarrow X$ , otherwise  $U \rightarrow R$

	IR	IX	R	RIX	U	X
IR	+	+	+	+	-	-
IX	+	+	-	-	-	-
R	+	-	+	-	-	-
RIX	+	-	-	-	-	-
U	-	-	+	-	-	-
X	-	-	-	-	-	-



## Multi-Granularity Locking (8)

### Intention Locks (contd.)



- Complete protocol (contd.)
  - ‚strict lock discipline‘ demanded
    - Lock requests from root to leaves
    - Before T requests R- or IR-lock for a node, it must hold IX- or IR-locks for all predecessors of this node
    - When a X-, U-, RIX- or IX-lock is requested all predecessor nodes must be hold in RIX- or IX-mode
    - Lock releases from leaves to root
    - At EOT all locks are to be released



## Multi-Granularity Locking (9)

### Intention Locks (contd.)


- Complete protocol (contd.)
  - Example
    - IR- and IX-Mode
    - TA1 reads  $t_1$  in T1
    - TA2 writes row in T2 (a)
    - TA3 reads T1 (b)



# Multi-Granularity Locking (10)

## Intention Locks (contd.)

- Complete protocol (contd.)
  - Example
    - RIX-Mode
      - TA1 reads all rows of T1 and updates  $t_3$
      - TA2 reads T1 (a)
      - TA3 reads  $t_2$  in T1 (b)




# CC with Versions (1)

## RAX

- Compatibilities

	R	A	X
R	+	⊕	-
A	⊕	-	-
X	-	-	-

- Example



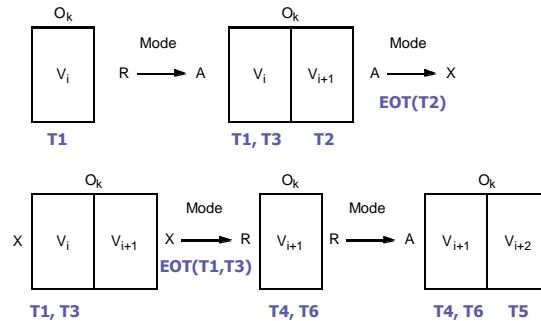
RAX: T2 → T1



## CC with Versions (2)

### ▪ RAX (contd.)

- Updates in temporary object copy




T4 (read),  
T5 (update),  
T6 (read)  
must wait, because of incompatibility of X



## CC with Versions (3)

### ▪ RAX (Forts.)

- Properties
  - Parallel read of current version is allowed
  - Writes are serialized as known (A-lock)
  - At EOT conversion of A- to X-locks, possibly wait for release of read locks (deadlocks may occur)
  - Higher concurrency than RX, but usually different serialization order:



## CC with Versions (4)

- **RAX (Forts.)**

- Problem
  - Not beneficial for mix of long read and short update transactions on shared objects
    - New version becomes available for new readers not before the old version has been abandoned
    - Severe obstructions of update transactions by (long) read transactions



## CC with Versions (5)

- **Multi-Version Concurrency Control**

- Idea
  - Update transactions create new object versions
    - Only one new version per object can be created
    - New version is released at EOT
  - Read transactions see the DB state which is valid at their BOT
    - They always access the youngest object version, which was released before their BOT
    - They do not acquire and adhere to locks
    - There is no blocking or aborts for read transactions; however, they possibly access older object versions






## CC with Versions (6)

### Multi-Version Concurrency Control (contd.)

- Example for Object  $O_k$

- Temporal order of accesses to  $O_k$


$T_j$ (BOT)	→	$V_i$ (current version)
$T_m(X)$	→	create $V_{i+1}$
$T_n(X)$	→	delay until $T_m$ (EOT)
$T_m$ (EOT)	→	release $V_{i+1}$
$T_n(X)$	→	create $V_{i+2}$
$T_j$ (Ref)	→	$V_i$
$T_n$ (EOT)	→	release $V_{i+2}$



## CC with Versions (7)

### Multi-Version Concurrency Control (contd.)

- Example



- Consequence


- Considerably less conflicts
  - Read TA are not taken into account by concurrency control
  - Update TA are synchronized (among each other) by a general concurrency control mechanism (locks, OCC, . . .)



## CC with Versions (8)

### Multi-Version Concurrency Control (contd.)

- Additional storage and maintenance overhead
  - Version pool management, garbage collection
  - Finding versions



- Storage optimization: versions at record level, compression techniques
- Application in some commercial DBMS (e.g. Oracle)

## Predicate Locks (1)

### Logical locks

- Locks by predicates (WHERE clause)
- Avoiding the phantom problem
- Elegant

### Form

- LOCK (R, P, a), UNLOCK (R, P)
  - R: Relation name
  - P: Predicate
  - $a \in \{\text{read, write}\}$
- Lock (R, P, write) locks all records of R (exclusively) which fulfill predicate P

Eswaran, K.P. et al.: The notions of consistency and predicate locks in a data base system. in: Comm. ACM 19:11, 1976, 624-633

## Predicate Locks (2)

- **Example:**

<b>T1:</b>	<b>LOCK(R1, P1, read)</b>	<b>T2:</b>	<b>...</b>
	<b>LOCK(R2, P2, write)</b>		<b>LOCK(R2, P3, write)</b>
	<b>LOCK(R1, P5, write)</b>		<b>LOCK(R1, P4, read)</b>
	<b>...</b>		<b>...</b>


- **Problem: conflict detection**

- General case recursively undecidable, even with restricted arithmetic operations
- Decidable class: simple predicates of the form  $(A \odot \text{Value}) \{\wedge, \vee\} (\dots)$



## Predicate Locks (3)

- **Decision procedure**



1. If  $R \neq R'$ , no conflict
2. If  $a = \text{read}$  and  $a' = \text{read}$ , no conflict
3. If  $P(t) \wedge P'(t) = \text{TRUE}$  for some  $t$ , then there is a conflict

- **Example**

<b>T1:</b>	<b>T2:</b>
<b>LOCK (Emp, Age &gt; 50, read)</b>	<b>LOCK (Emp, Emp-Id = 4711, write)</b>



## Predicate Locks (4)

### ▪ Drawbacks

- Costly decision procedure; generally many predicates ( $N > 100$ )
- pessimistic decision  $\rightarrow$  restriction of parallelism
- For descriptive languages only!
- Special case:  $P=\text{TRUE}$  is equivalent to relation lock  $\rightarrow$  large lock granulates, low parallelism



## Predicate Locks (5)

### ▪ More Efficient Implementation: Precision Locks

- Predicate locks only for read data
- Write locks for updated rows
- No (more) need to test, whether or not two predicates are disjunct
- Easier test, whether or not row fulfills predicate
- Data structures
  - Predicate list: read locks of current TAs are described by predicates  
(Emp: Age > 50 and Occupation = 'Prog.')  
(Emp: Name = 'Meier' and Salary > 50000)  
(Abt: Anr = K55)  
...
  - Update list: contains updated records of current TAs  
(Emp: 4711, 'Müller', 30, 'Prog.', 70000)  
(Dept: K51, 'DBS', ...)  
...

J.R. Jordan, J. Banerjee, R.B. Batman: Precision Locks, in: Proc. ACM SIGMOD, 1981, 143-147



## Predicate Locks (6)

### ▪ Precision locks (contd.)

- Read request (predicate P):
  - For each record in the update list it has to be tested, whether or not it fulfills P
  - If so → conflict
- Write request (tuple T):
  - For each predicate p in predicate list P(T) is to be checked
  - If T does not fulfill a predicate → write lock is granted



## „High-Traffic“-Synchronization (1)

### ▪ Synchronization of High-Traffic-Objects

- High-Traffic-Objects
  - Mostly numerical fields with aggregated information, e.g.
    - Number of free seats
    - Sum of bank balances
  - Problem: locks result in high number of waiting transactions
- Treatment
  - Solution of lock problems:
    - Either: avoiding such fields during modeling
    - or: optimistic or not-locking concurrency control mechanisms



## ,High-Traffic'-Synchronization (2)

### ▪ Escrow Approach

- Declaration of high-traffic objects as Escrow Fields
- Special operations
  - Reservation of a certain set of values w.r.t to field F1

```
IF      ESCROW (field=F1, quantity=C1, test=(condition))
THEN    'continue with normal processing'
ELSE    'perform exception handling'
```
  - Update of Escrow Field without synchronization

```
USE (field=F1, quantity=C1)
```
- Reservation of Escrow Field only if (optional) test condition fulfilled
- If reservation successful then it can not be invalidated afterwards

P. O'Neil: The Escrow Transactional Method,  
in: ACM Trans. on Database Systems 11: 4, 1986, 405-430



## ,High-Traffic'-Synchronization (3)

### ▪ Escrow Approach (contd.)

- Current value of Escrow Field
  - unknown, if there are reservations of running TAs
  - Value interval that contains all possible values
  - For value  $Q_k$  of Escrow Field  $k$  the following is true
    - $LO_k \leq INF_k \leq Q_k \leq SUP_k \leq HI_k$
    - Adjustment of INF, Q, SUP at request, Commit and Abort of TA
  - Properties
    - Range test w.r.t to value interval
    - Minimal-/maximal values (LO, HI) must not be exceeded
    - High parallelism possible
  - Drawbacks
    - Special programming interface
    - Real value not known



## „High-Traffic“-Synchronization (4)

### ▪ Escrow Approach (contd.)



- Current value of Escrow Field (contd.)
  - Example: Field with LO = 0, HI = 500 (number of free seats)

T1	T2	T3	T4	INF	Q	SUP
				15	15	15
-5				10	10	15
	-8			2	2	15
		+4		2	6	19
			-3	-1		
c				2	6	14
		c		6	6	14
	a			14	14	14

## Analysis of Lock Mechanisms (1)

### ▪ Synchronization and Load Control

- Load Control
  - „blind“ throughput maximization
    - More active TA → more locked objects → higher conflict risk → longer waiting periods, more deadlocks → even more active TA
  - Multi Programming Level (MPL)
    - Determines performance, number of conflicts and aborts
    - Danger of „Thrashing“ if critical MPL-Value is exceeded



## Analysis of Lock Mechanisms (2)

### ▪ Synchronization and Load Control (contd.)

- Dynamic Load Control
  - „Static“ MPL adjustment not appropriate
    - Changing load properties, multiple transaction types
  - Idea:
    - Dynamic MPL adjustment in order to avoid „Thrashing“
  - Approach: Conflict Rate of locking mechanisms
    - Conflict Rate =  $\frac{\# \text{ hold locks}}{\# \text{ locks of non-blocked transactions}}$
    - Critical value:  $\sim 1,3$  (determined empirically)
    - New transactions only if critical value not reached yet
    - Abort of transactions if critical value is exceeded

Weikum, G. et al.: The Comfort Automatic Tuning Project,  
in: Information Systems 19:5, 1994, 381-432



## Conclusion (1)

### ▪ Concurrency Control by Locking

- Locks ensure that history stays serializable
- As soon as a conflict operation is submitted object access is blocked
- Multiple variants (of locking mechanisms)
- Predicate locks represent an elegant idea, but are not feasible for practical use; possibly exploitation in the form of precision locks
- DBS-Standard: multiple lock granulates by hierarchical lock mechanisms
- Locking is pessimistic and universally applicable
- Deadlocks are inherent problem of locking/blocking mechanisms





## Conclusion (2)

- **Further Mechanisms**

- RAX limits number of versions and reduces blocking periods for certain situations
- Multi-version mechanisms deliver high degree of parallelism and less deadlocks; however, they cause higher overhead (algorithm, storage, ...)
- Simple OCC- (and timestamp) mechanisms cause too many aborts



## Conclusion (3)

- **„Hard“ Concurrency Control Problems**

- Examples
  - 'Hot Spots' / 'High Traffic'-Objects
  - long (Update-) TA
- Solutions
  - If there is no avoidance then special protocols required
  - Exploitation of semantic knowledge about operations / objects in order to reduce number of conflicts
- however
  - (Possibly) extension of programming interface
  - Restricted applicability
  - Additional overhead
- Dynamic Load Control required
  - Avoidance of „Thrashing“ in case of changing load situations, multiple transaction types, ...
  - Consideration of Conflict Rate ( $\sim 1.3$ ) for dynamic MPL adjustment

