

Chapter 5

Logging and Recovery

Requirements / Basic Notions,
Logging,
Insertion Strategy (Non-/Atomic),
Propagation Strategy (No-/Force),
Replacement Strategy (No-/Steal),
Save Points,
Recovery: Restart-Procedure



Requirements / Basic Notions (1)

- **DBMS Task**
 - Automatic Handling of expectable failures
- **Expectable Failures**
 - DB-Operation rejected
 - Commit not accepted
 - Power breakdown
 - Devices do not work (e.g. magnetic disk)
 - ...
- **Special Characteristics of DBMS Failure Handling**
 - Restriction to and reparation of runtime failures (failure tolerant systems)
 - „Reparation“ of static DB structures




Requirements / Basic Notions (2)

- **General Problems**
 - Failure detection
 - Failure localization
 - Estimation of damage
 - Recovery (itself)
- **Failure Model of centralized DBMS**
 - Transaction failure
 - System failure
 - Device failure
 - Disaster
- **Precondition**
 - Collecting redundant information during normal operation (Logging)



Requirements / Basic Notions (3)

- **Transaction paradigm requires**
 - "All or Nothing" (Atomicity)
 - Durability
- **Goal of Recovery**
 - Most recent transaction-consistent DB state
- **System Environment?**
 - Operating system, application system, other components



Requirements / Basic Notions (4)

▪ Basic Forms of Recovery

- Forward-Recovery
 - Find a state, at which system can continue to operate
 - However, non-stop paradigm not generally applicable
- Backward-Recovery
 - Back to most recent consistent state and further processing from there
 - Requires that at all abstractional layers it is clearly defined, to which state it must be restored in case of failure



Requirements / Basic Notions (5)

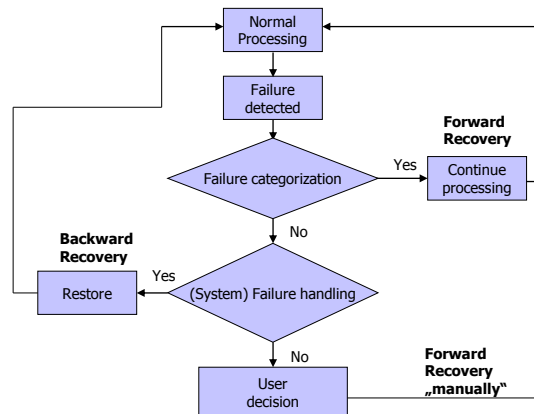
▪ Overhead

- "A recoverable action is 30% harder and requires 20% more code than a non-recoverable action" (J. Gray)
- Statement and transaction atomicity required
 - 2 principles
 - **Do things twice:**
prepare first; if OK then concrete modification
 - **Do things once:**
immediate modification; in case of failure internal restore
 - Usually second principle is used (more optimistic and efficient)



Requirements / Basic Notions (6)

Basic processing



Requirements / Basic Notions (7)

Failure classes

Effect	Failure type	Failure class
A single transaction	Violation of system restrictions Security violation Excessive resource requisition Application failures e.g. wrong operations or values	<i>Transaction failure</i>
Several transactions	Planned system shut down Problems of resource allocation System overload Deadlock	<i>System failure</i>
All transactions (overall system)	System break down with loss of main memory contents Damage of secondary storage Damage of computer center	<i>System failure</i> <i>Device failure</i> <i>Disaster</i>

Requirements / Basic Notions (8)

- **Preconditions for Recovery**

- quasi-stable storage
- accurate DBMS code
- accurate Log data
- independence of failures

- **Recovery Classes**

- 1. **Transaction Recovery** (R1)

- UNDO of single not-committed transaction during database operation (transaction failure, deadlock)
 - Forms
 - Complete UNDO to initial state
 - Partial UNDO to savepoint within transaction



Requirements / Basic Notions (9)

- **Recovery Classes (contd.)**

- 2. **Crash Recovery** (R2) after System Crash

- Restore of most recent transaction consistent DB state
 - Necessary actions
 - (partial) REDO of successful transactions (REDO of lost modifications)
 - UNDO of all interrupted transactions (removal of all their modifications from permanent DB)

- 3. **Media Recovery** (R3) after Device Failure

- Mirroring (at disk level)
 - Complete REDO of all modifications of successful completed transactions on archive copy of DB



Requirements / Basic Notions (10)

- **Recovery Classes (contd.)**

- 4. **Disaster Recovery** (R4)

- DB copy in remote system
 - Delayed continuation of DB processing on repaired/new system on basis of archive copy (possibly data loss)



Requirements / Basic Notions (11)


- **Not (formally) classified**

- R5 Recovery
 - Log-Data damaged
 - R6 Recovery
 - Beyond DBMS
 - Compensation transactions
 - Manual treatment



Requirements / Basic Notions (12)

■ DB-Recovery – System Components



Requirements / Basic Notions (13)

■ DB-Recovery – System Components (contd.)

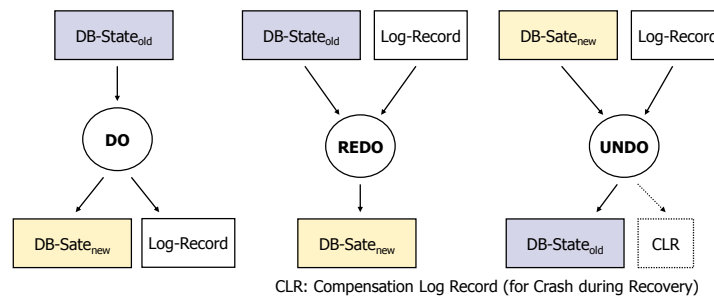
- Buffering Log Data in Main Memory (Log Buffer)
 - Propagation (at the latest) at Commit
- Usage of Log Data
 - Temporary Log File for Handling Transaction Failures and System Failures
 - DB + temp. Log \Rightarrow DB
 - Handling Device Failures
 - Archive Copy + Archive Log \Rightarrow DB

Logging (1)

■ Task


- Collecting redundant data w.r.t. to modifications during normal DB processing
- Usage in case of failure (Undo-, Redo-Recovery)

■ Do-Redo-Undo-Principle



Logging (2)


- Logging can be performed at (DBMS-) System Layer
- Logging Techniques:



Logging (3)

■ Logical Logging


- Logging update DML operations with parameters
- General Problem
 - Set oriented update operations
- Precondition
 - After system crash persistent data base must at least be action consistent, in order to being able to perform ,reverse operations'
- Deferred (indirect) insertion strategy needed



Logging (4)

■ Physical Logging


- Log Granulate: Page vs. Entry/Record
 - State Logging
 - Before Images and After-Images are stored in Log File
 - Transition Logging
 - Difference between Before- and After-Image is stored
 - Applicable with direct as well as indirect insertion strategies
- ### ■ Problems of logical and physical Logging Techniques
- Logical Logging: not compatible to Update-in-Place
 - Physical, „byte-oriented“ Logging: complex and inflexible w.r.t. deletion and insertion operations



Logging (5)

■ Physiological Logging

- Combined physical/logical Logging
 - physical-to-a-page, logical-within-a-page
 - Each Log-Record relates to a single database page
 - Logging of elementary, internal operations within a page
 - Compatible to Update-in-Place



Logging (6)

■ Examples

- Modifications of a page A
 1. Inserting object a into page A ($A_1 \rightarrow A_2$)
 2. Modifying object b_{old} to b_{new} in page A ($A_2 \rightarrow A_3$)

	<i>logical</i>	<i>physical</i>
<i>States</i>		Logging Before- and After-Images 1. A_1 and A_2 2. A_2 and A_3
<i>Transitions</i>	Logging Operations with Parameters 1. Insert (a) 2. Update (b_{old}, b_{new})	Logging 'Differences' 1. $A_1 \oplus A_2$ 2. $A_2 \oplus A_3$

Logging (7)

Examples (contd.)

- Logging ,Diffs': Reconstruction of Pages
 - A1 as starting point or A3 as endpoint available
 - REDO-Recovery
 - $A_1 \oplus (A_1 \oplus A_2) = A_2$
 - $A_2 \oplus (A_2 \oplus A_3) = A_3$
 - UNDO-Recovery
 - $A_3 \oplus (A_2 \oplus A_3) = A_2$
 - $A_2 \oplus (A_1 \oplus A_2) = A_1$

A	B	XOR
0	0	0
0	1	1
1	0	1
1	1	0

Logging (8)

Assessment of Logging Techniques

	Overhead during normal processing	Restart-Overhead after failure (Crash)
Page-Logging	--	+
Page-Transition-Logging (Differences)	-	+
Entry/Record-Logging / physiological Logging	+	+
Logical Logging	++	--

Logging (9)

■ Entry-Logging vs. Page-Logging

- Advantages of Entry-Logging
 - Low storage overhead
 - Less Log-E/As
 - Allows 'better' buffering of Log data (Group-Commit)
 - Supports more fine-grained concurrency control granulates (Page-Logging → CC on page level)
- Drawback of Entry-Logging
 - Recovery more complex than in case of Page-Logging
 - e.g. before application of log records pages must be loaded into main memory



Logging (10)

■ Structure of (temporary) Log-File

- Several different Log Records needed
 - BOT-, Commit-, Abort-Records
 - Update-Record (UNDO-Information, e. g. 'Before-Images', and REDO-Information, e. g. 'After-Images')
 - Checkpoint-Records
- Update Record
 - Structure of Record:
[LSN, TAID, PageID, Redo, Undo, PrevLSN]
 - LSN: Log Sequence Number
 - Unique ID of log record
 - LSNs are created in monotonically ascending order
 - Thus, chronological order of logging entries can be reconstructed



Logging (11)

- **Structure of (temporary) Log-File (contd.)**
 - ID of transaction which issued update
 - PageID
 - ID of modified page
 - If modification relates to more than one page then several log records need to be created
 - Redo
 - Redo-Information specifies how modification can be reproduced
 - Undo
 - Undo-Information specifies how modification can be withdrawn
 - PrevLSN
 - Pointer to previous log record of the same transaction
 - Needed for efficiency reasons during transaction UNDO

Logging (12)

- **Structure of (temporary) Log-File (contd.)**
 - Example

Schritt	T ₁	T ₂	Log
			[LSN, TAID, PageID, Redo, Undo, PrevLSN]
1.	BOT		[#1, T ₁ , BOT , 0]
2.	r(A, a ₁)		
3.		BOT	[#2, T ₂ , BOT , 0]
4.		r(C, c ₂)	
5.	a ₁ := a ₁ - 50		
6.	w(A, a ₁)		[#3, T ₁ , P _A , A+=50, A+=50, #1]
7.		c ₂ := c ₂ + 100	
8.		w(C, c ₂)	[#4, T ₂ , P _C , C+=100, C+=100, #2]
9.	r(B, b ₁)		
10.	b ₁ := b ₁ + 50		
11.	w(B, b ₁)		[#5, T ₁ , P _B , B+=50, B+=50, #3]
12.	Commit		[#6, T ₁ , Commit , #5]
13.		r(A, a ₂)	
14.		a ₂ := a ₂ - 100	
15.		w(A, a ₂)	[#7, T ₂ , P _A , A-=100, A-=100, #4]
16.		Commit	[#8, T ₂ , Commit , #7]

Logging (13)

■ Structure of (temporary) Log-File (contd.)

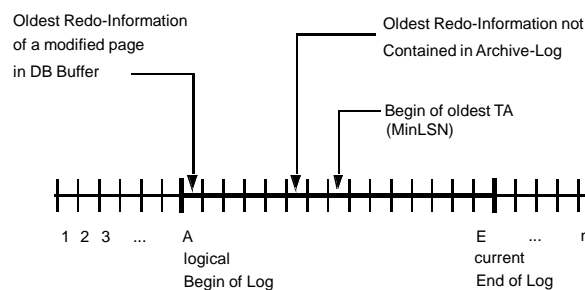
- Sequential File
 - Writing new logging data at end of file
- Log-Data relevant for Crash Recovery only for restricted period of time
 - Undo-Records no longer needed as soon as transaction is completed successfully
 - After insertion of page into permanent DB Redo-Information is no longer needed
 - Redo-Information for Media-Recovery is collected in Archive-Log!



Logging (14)

■ Structure of (temporary) Log-File (contd.)

- Ring Buffer



Related System Components (1)

▪ Overview

- Insertion Strategy
 - Direct insertion of modifications into permanent DB (*non-atomic*)
 - Deferred insertion (*atomic*)
- Replacement Strategy
 - Pushing 'dirty' pages to secondary storage (*steal*)
 - Only pages of successfully completed transactions are pushed (*no steal*)
- Propagation Strategy
 - Propagation at Commit mandatory (*force*)
 - Propagation possibly after Commit (*no force*)
- Locking Granulate
- Commit-Procedure



Related System Components (2)

▪ Insertion Strategy

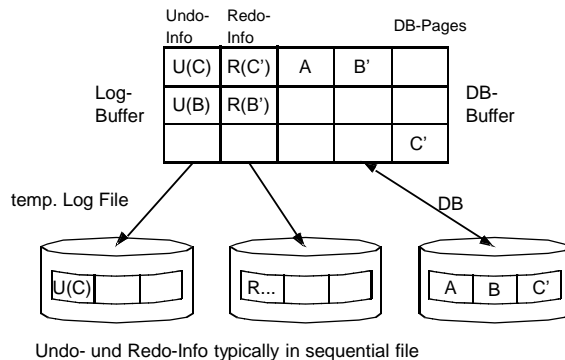
- *non-atomic / direct / update-in-place*
 - Modified page is always stored back to the same block on disk
 - ‚storing back‘ (here) means insertion into permanent DB
 - ‚atomic‘ insertion of several pages not possible (non-atomic)
- Requirements
 - WAL-Principle: *Write Ahead Log* for Undo-Info;
U(B) before B' (cf. next slides)
 - Logging Redo-Info at the latest at Commit;
R(C') + R(B') before Commit (cf. next slides)



Related System Components (3)

■ Insertion Strategy (contd.)

- *non-atomic / direct / update-in-place (contd.)*



Related System Components (4)


■ Insertion Strategy (contd.)

- *atomic / deferred*
 - e. g. in System R, SQL/DS
 - Modified page is written to separate block on disk, actual insertion into DB' is performed later on
 - Page table contains page address
 - Deferred, atomic insertion of multiple modifications possible by switching page tables
 - Action consistent or even transaction consistent DB on disk
 - Thus, logical logging applicable
 - Requirements
 - WAL-Principle:
U(C) + U(B) before checkpoint
 - R(C') + R(B') at the latest at Commit

Related System Components (5)

■ Insertion Strategy (contd.)

- *atomic / deferred (contd.)*



Related System Components (6)

■ Replacement Strategy

- Problem: Replacement of 'dirty' pages
- *steal*
 - Modified pages can be replaced (in buffer) and stored into the permanent DB at any time, esp. before commit of corresponding TA
 - Higher flexibility for page replacement
 - Undo-Recovery needed (TA Abort, System Crash)
 - *steal* requires observation of WAL principle, i.e., before writing a dirty page corresponding UNDO Information (e.g. Before Image) must be written into Log File
- *no steal*
 - Dirty pages must not be replaced
 - No UNDO Recovery needed
 - Problems in case of long Update TA

Related System Components (7)

▪ Propagation Strategy

- *force*
 - All modified pages are propagated to permanent DB at the latest at Commit ('writing through')
 - No Redo Recovery needed after System Crash
 - High Overhead
 - Large DB Buffers possibly not really exploited
 - Longer answering times for update TA
- *noforce*
 - No 'write through' at Commit
 - At (the latest at) Commit only Redo Information must be written to Log File
 - Redo Recovery after System Crash
- Commit-Rule
 - before TA Commit sufficient Redo Information (e. g. *After Images*) must be written for all modifications



Related System Components (8)

▪ Consequences


	steal	nosteal
force	UNDO NO REDO	NO UNDO NO REDO
noforce	UNDO REDO	NO UNDO REDO



Related System Components (9)

Lock Management


- Log Granulate must be smaller as or equal to Log Granulate
 - Example
 - Record Locks
 - Before* and *After Images* at page level
 - Undo (Redo) of a modification can overwrite parallel modification of the same page (*lost update*)



Related System Components (10)

Commit Procedure

- Requirements
 - Modifications need to be 'assured' at Commit
 - Modifications get visible to other TA not before it can be assured that corresponding update TA will get to its Commit (Problem of recursive Aborts)



Related System Components (11)

▪ Commit Procedure (contd.)

- 2-phase processing
 - Phase 1: ensuring repeatability of TA
 - Storing modifications
 - Writing Commit Record to Log
 - Phase 2: making modifications visible to others (Releasing Locks)
 - At the end of phase 1 user/application can be informed that TA has been successful
- Example: Commit Procedure for *force, steal*:
 1. Writing Before Images to Log
 2. Force of modified DB Pages
 3. Writing After-Images (for Archive Log) Commit RecordIn case of NoForce just 3. needed for first Commit Phase



Related System Components (12)

▪ Commit Procedure (contd.)

- Group Commit
 - Log File is potential bottleneck
 - At least 1 Log-E/A for each update TA
 - max. about 250 sequential writes per second (1 disk)
 - Group Commit means writing Log Data of several TA
 - Buffering Log Data in Log Buffer (1 or more pages)
 - Precondition: Record Logging
 - Log Buffer is written to Log File if full or time limit exceeded (Timer)
 - Insignificant delay of Commit



Related System Components (13)

■ Commit Procedure (contd.)


- Group Commit (contd.)
 - Group Commit allows reduction to 0.1 - 0.2 Log-I/Os per TA
 - Less CPU overhead (for I/Os) reduces waiting times for CPU
 - Dynamic adjustment of timer value by DBMS desirable
 - Thus, Group Commit allows increase of throughput, esp. w.r.t. log bottleneck and high CPU utilization




Related System Components (14)

■ Commit Procedure (contd.)

- Comparison
 - Standard 2PC




- Group Commit



Related System Components (15)

Commit Procedure (contd.)

- Further Optimization: *Pre-Commit*
 - Lock release after Commit Record has been written to Log Buffer (not Log File)
 - TA can only be aborted by system crash
 - In this case all depending TAs (having seen 'unstored' modifications because of early lock release) fail too




- In all 3 variants user is informed about TA termination not before Commit Record has been written

Checkpoints (1)

Motivation: Checkpoints

- Goal: restricting REDO overhead after System Crash (noforce)
- Without checkpoint potentially all modifications since system start would have to be 'repeated'
- Especially critical: Hot-Spot-Pages



Checkpoints (2)

■ Management Data

- Log File
 - BEGIN_CHKPT Record
 - (actual) Checkpoint Information, e. g. list of active TAs
 - END_CHKPT Record
- Log Address of last Checkpoint Record is kept in special System File

■ Checkpoints and Insertion`


- *atomic*
 - State of permanent DB is state of last (successful) checkpoint
- *non-atomic*
 - State of permanent DB contains all pages inserted before crash



Checkpoints (3)

■ Kinds of Checkpoints


- Direct Checkpoints
 - All modified pages are stored into the permanent DB
 - REDO Recovery starts at last Checkpoint
 - Drawback: long system down time, because there must be no modifications ,during Checkpoint`
 - *Transaction consistent or action consistent Checkpoints*



Checkpoints (4)

■ Kinds of Checkpoints (contd.)



- Indirect/Fuzzy Checkpoints
 - Modified pages do not have to be propagated
 - Only state information (IDs of pages in Buffer, active TAs, open files etc.) are written to Log File
 - Minor (Checkpoint-) Overhead
 - Generally, REDO Information before checkpoint must be taken into account
 - Special treatment of Hot-Spot-Pages



Checkpoints (5)

■ Transaction Oriented Checkpoints


- Force can be considered as special Checkpoint: only pages of a single TA are written
- Checkpoint always relates to exactly one TA (TOC = transaction-oriented checkpoint \equiv *force*)



Checkpoints (6)

Transaction Oriented Checkpoints (contd.)

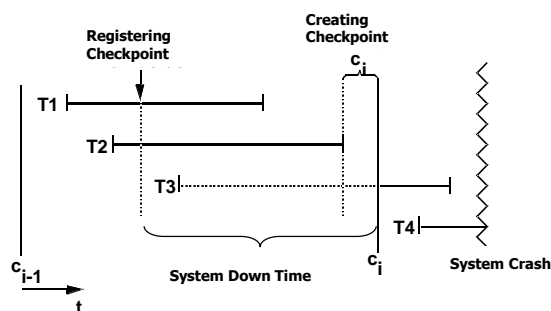

- Properties
 - Commit processing enforces propagation of all modified pages of the TA
 - Insertion of all modifications into permanent DB
 - Comment in Log File
 - only *atomic* supports atomic insertion of multiple pages
 - Thus, at least in case of direct Insertion of pages UNDO Recovery has to be provided (*steal*)



Checkpoints (7)

Transaction Consistent Checkpoints


- Checkpoint always relates to all TA
(TCC = transaction consistent checkpoint)



Checkpoints (8)

Transaction Consistent Checkpoints (contd.)

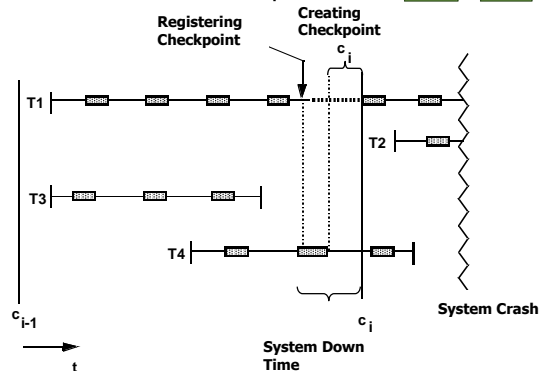

- Properties
 - Propagation to be deferred until end of all active update TA
 - New update TA must wait, until checkpoint creation completed
 - Crash Recovery starts at last checkpoint



Checkpoints (9)

Action Consistent Checkpoints


- Checkpoint always relates to all TA
(ACC = action consistent checkpoint)



Checkpoints (10)

▪ Action Consistent Checkpoints (contd.)


- Properties
 - No update statements during checkpoint
 - In comparison to TCC shorter down times, but lower checkpoint ‚quality‘
 - Crash Recovery not limited by last checkpoint



Checkpoints (11)



▪ Fuzzy Checkpoints

- DB stays ‚fuzzy‘
 - Only relevant for Update-in-Place (*non-atomic*)
- Problem
 - Determination of Log position, at which Redo Recovery has to start
 - Buffer manager stores StartLSN for each modified page, i.e., LSN of first modification since reading from disk
 - Redo-Recovery after Crash starts at MinDirtyPageLSN (= MIN(StartLSN))
- Checkpoint Information
 - MinDirtyPageLSN, List of active TA and their StartLSNs, ...



Checkpoints (12)

Fuzzy Checkpoints (contd.)

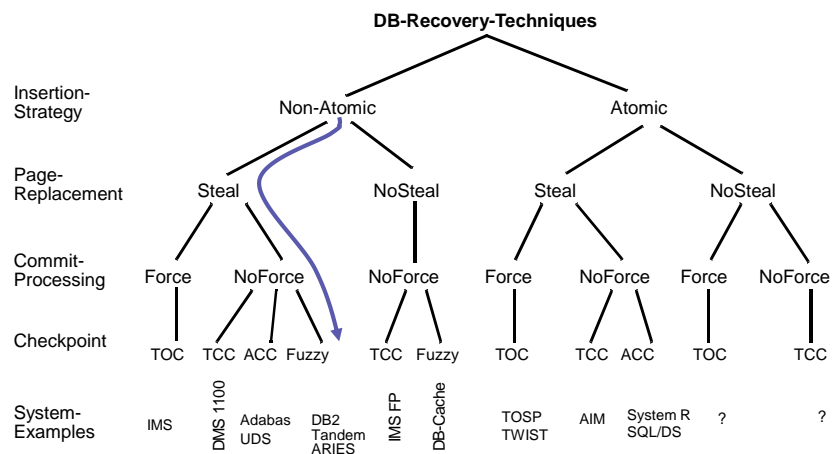


- Modified pages are written asynchronously
 - If necessary, creating copy of page (for Hot-Spot-Pages)
 - Propagating page
 - Adjusting StartLSN



Checkpoints (13)

Combinations



Recovery (1)

■ LSNs


- Challenge
 - Decision at restart, whether or not Recovery action is necessary for considered page (old or new state on external storage?)
 - For that purpose each page header contains LSN of most recent Log Entry L related to this page (PageLSN (B) := LSN (L))
- Decision Procedure
 - Restart contains Redo- and Undo-Phase
 - Redo necessary, only if
 $\text{Page-LSN} < \text{LSN of Redo-Log-Entry}$
 - Undo necessary, only if
 $\text{Page-LSN} \geq \text{LSN of Undo-Log-Entry}$



Recovery (2)

■ LSNs (contd.)

- Simplified application: Page Locks




- Redo T1: $S(10) = T1(10): -$
 $S(10) < T1(20): \text{Redo, } S(20)$
- Undo T2: $S(20) < T2(40): -$
- Page-LSN is updated at Redo (increases monotonically)



Recovery (3)

LSNs (contd.)

- Simplified application: Page Locks (contd.)



- Redo T1: $S(50) > T1(10):$ –
 $S(50) > T1(20):$ –
- Undo T2: $S(50) < T2(60):$ –
 $S(50) \geq T2(50):$ Undo
 $S(50) \geq T2(40):$ Undo



Recovery (4)

LSNs (contd.)

- Undo in LIFO Order
 - Special treatment of 'UNDOS' necessary so that repeated application leads to the same result (idempotent)
 - State Logging and LIFO Order ensure that processing is idempotent



Recovery (5)

■ Crash-Recovery

- Goal
 - Creating the most recent transaction consistent DB state from permanent DB and temporary Log File
- In case of Update-in-Place (*non-atomic*)
 - State of permanent DB after Crash unpredictable („chaotic“)
 - Thus, only physical (or physiological) Logging applicable
 - A Block of the permanent DB either is
 - Up-to-date
 - or outdated (*noforce*) → Redo
 - or ‚dirty‘ (*steal*) → Undo



Recovery (6)

■ Crash-Recovery (Forts.)

- In case of *atomic*
 - State of permanent DB corresponds to the most recent successful propagation (checkpoint)
 - At least action consistent → DML statements can be executed (logical Logging)
 - *force*: no Redo
 - *noforce*:
 - Transaction consistent propagation → Redo, no Undo
 - Action consistent propagation → Undo + Redo



Recovery (7)

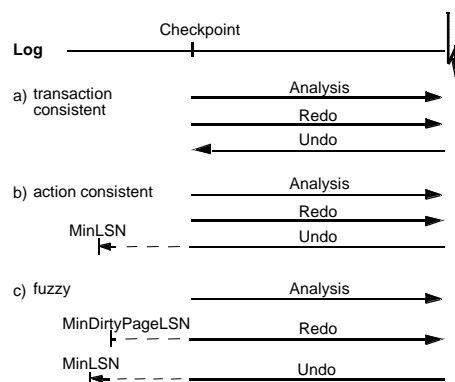
General Restart Procedure

- 3 Phases
 1. Analysis Phase
 - From last checkpoint to end of Log
 - Determination of winner and loser TAs as well as of modified pages
 2. Redo Phase
 - Reading Log forward: starting point depends on checkpoint type
 - selective Redo (*redo winners*) in case of page locks or complete Redo (*repeating history*)
 3. Undo Phase
 - UNDO of all 'losers'
 - Reading Log backward until BOT record of oldest loser TA



Recovery (8)

General Restart Procedure (contd.)



Recovery (9)


Restart Procedure (Update-in-Place)

- Properties: non-atomic, steal, noforce, fuzzy checkpoints
- Process
 1. Analysis Phase
 - From last checkpoint to end of log
 2. Redo Phase
 - Starting point depends on checkpoint type: here MinDirtyPageLSN
 - selective Redo: modifications of winner TAs only
 3. Undo Phase
 - Loser TA up to MinLSN



Recovery (10)

Restart Procedure (Update-in-Place) (contd.)



- Overhead
 - For steps 2 and 3 corresponding pages must be loaded from external storage
 - Page LSNs indicate, whether or not Log information must be applied
 - At the end all modified pages must be propagated again, or a checkpoint is created, respectively



Recovery (11)

Redo

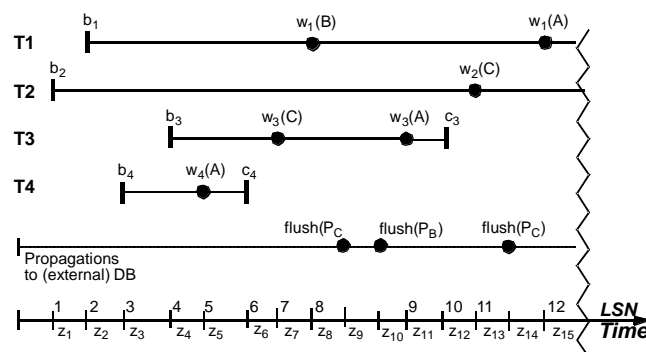
- In case of physical and physiological Logging
 - Redo action for Log record L is determined by PageLSN of corresponding page B


```
if (B not in Buffer) then
    load B into main memory;
fi;
if (LSN (L) > PageLSN (B)) then
    Redo (modification from L);
    PageLSN (B) := LSN (L);
fi;
```
- Repeated application of Log record (e.g. after multiple failures) keeps correctness (REDO is idempotent)
- Recovery in case of Crash during Restart?



Recovery (12)

Restart Example:



Recovery (13)

Restart Example (contd.)

- Assumptions
 - At the beginning: all Page-LSNs 0
- Analysis Phase:
 - Winner-TA: T3, T4
 - Loser-TA: T1, T2
 - relevant pages: P_A, P_B, P_C
- Comment
 - In the example: page locks
 - Thus, selective Redo sufficient (Redo only for winners)



Recovery (14)

Restart Example (contd.)

- Redo Phase:
 - Checking Log records for T₃ and T₄ (forward)

TA	Page	Page-LSN	Log-Record-LSN	Action
T ₄	P _A	0 → 5	5	REDO
T ₃	P _C	11	7	no REDO
T ₃	P _A	5 → 9	9	REDO

- Redo only, if Page-LSN < Log-Record-LSN
- Page-LSNs increase monotonically



Recovery (15)

Restart Example (contd.)

- Undo Phase:
 - Checking Log records for T_1 and T_2 (backward)

TA	Page	Page-LSN	Log-Record-LSN	Action
T_1	P_A	9	12	no Undo
T_2	P_C	11	11	Undo
T_1	P_B	8	8	Undo


- Undo only, if Page-LSN \geq Log-Record-LSN
- Because of page logs there is no interference between REDO and UNDO actions; state logging ensures that UNDO is idempotent



Recovery (16)

UNDO Problems w.r.t. LSN Exploitation

- Problem 1: TA UNDO
 - Taking previous Rollbacks into account?



- Redo of T2: $S(10) < T2(40)$: Redo, $S(40)$
- Undo of T1: $S(40) > T1(20)$: Undo, Failure



Recovery (17)

■ UNDO Problems w.r.t. LSN Exploitation (contd.)


- Problem 1: TA UNDO (contd.)
 - Comment
 - UNDO of modification 20, although modification not represented by page S
 - Assigning LSN = 20 to S violates monotonicity requirement for Page LSNs



Recovery (18)

■ UNDO Problems w.r.t. LSN Exploitation (contd.)

- Problem 2: Record Locks
 - T1 and T2 modify page S concurrently



- Redo T1:

$S(10) \geq T1(10):$	no Redo
$S(10) < T1(30):$	Redo, S(30)
- Undo T2 (LIFO):

$S(30) < T2(50):$	no Undo
$S(30) > T2(20):$	Undo, Failure!
- More general UNDO processing needed!



Recovery (19)

■ Failure tolerance of Restart

- Requirement: Restart must be idempotent

$$\text{Undo}(\text{Undo}(\dots(\text{Undo}(A))\dots)) = \text{Undo}(A)$$
$$\text{Redo}(\text{Redo}(\dots(\text{Redo}(A))\dots)) = \text{Redo}(A)$$

- Solution


- REDO idempotent since Page LSNs increase monotonically
- ‚Compensation Log Records‘ for UNDO




Recovery (21)

■ Failure tolerance of Restart

- CLR = Compensation Log Record
 - CLR for all Undo actions: Rollback and Undo Phase
 - in Redo Phase: complete Redo of Winners and Losers („repeating history“)
 - Illustration of Log File



- After complete Restart:



Recovery (22)

■ Failure tolerance of Restart (contd.)


- CLR = Compensation Log Record (contd.)
 - Redo Information of CLR equates to UNDO operation as performed in UNDO phase
 - CLRs are needed for repeated Restart (Crash during Restart); then their REDO Information is applied and corresponding Page LSNs are modified → idempotent
 - CLRs do not need Undo Information; they are skipped in subsequent Undo Phases (UndoNxtLSN)



Recovery (23)

■ Compensation Log Records (contd.)

- Example
 - all modifications relate to page S
 - State after Crash 1



Recovery (24)

■ Compensation Log Records (contd.)


- Example (contd.)
 - State after Crash 1 (contd.)
 - repeating history:
 - $S(40) > T1(10)$: –
 - ...
 - $S(40) \geq T2(40)$: –
 - $S(40) < T1(50)$: Redo, $S(50)$
 - $S(50) < T1(60)$: Redo, $S(60)$
 - $S(60) < T2(70)$: Redo, $S(70)$
 - Undo of T2:
 - CLR(80): Compensation of $T2(70)$, $S(80)$
 - Propagating S to DB (Flush S)
 - CLR(90): Compensation of $T2(40)$, $S(90)$
 - Crash



Recovery (25)

■ Compensation Log Records (contd.)

- Example (contd.)
 - State after Crash 2



Recovery (26)

■ Compensation Log Records (contd.)

- Example (contd.)
 - State after Crash 2 (contd.)
 - Repeating History:
S(80) > T1(10): –
...
S(80) > T2(70): –
CLR(80): –
CRL(90): Compensation of T2(40), S(90)
 - Undo of T2:
CLR(100): Compensation of T2(20), S(100)
 - End



Recovery (27)

■ Restart Procedure (Update-in-Place)


- Properties: non-atomic, steal, noforce, fuzzy checkpoints
 1. Analysis Phase
 - From last checkpoint to end of log
 2. Redo Phase
 - Starting point: MinDirtyPageLSN
 - Selective Redo or Repeating History (if necessary)
 3. Undo Phase
 - UNDO of losers back to MinLSN



Recovery (28)

Restart Procedure (Update-in-Place) (contd.)

- Properties: non-atomic, steal, noforce, fuzzy checkpoints (contd.)



- ARIES


- Algorithm for Recovery and Isolation Exploiting Semantics
- Developed by C. Mohan et al. (IBM Almaden Research)
- Realized in several commercial DBMS

Mohan, C. et al.: A Transaction Recovery Method Supporting Fine-Granularity Locking and Partial Rollbacks Using Write-Ahead Logging, in ACM TODS 17:1, 1992, 94-162



Recovery (29)

Restart Example 2



Recovery (30)

Restart Example 2 (contd.)

- Analysis Phase
 - Winner-TA: T_1
Loser-TA: T_2
relevant Pages: P_A, P_B, P_C
- Comment
 - Complete Redo
 - For Undo operations: CLR with the following structure:

[LSN, TAID, PageID, Redo, PrevLSN, UndoNextLSN]



Recovery (31)

Restart Example 2 (contd.)

- Redo Phase
 - Checking log records of all TA (T_1, T_2) forwards

TA	Page	Page-LSN	Log-Record-LSN	Action
T_1	P_A	7	3	No REDO
T_2	P_C	$0 \rightarrow 4$	4	REDO
T_1	P_B	$0 \rightarrow 5$	5	REDO
T_2	P_A	7	7	No REDO

- Redo, if Page-LSN < Log-Record-LSN



Recovery (32)

Restart Example 2 (contd.)

- Undo Phase
 - Checking log records of loser TA T_2 backward
 - For each log record Undo is performed and CLR written to log end

TA	Log-Record-LSN	Action
T_2	7	UNDO and CLR[8, T_2 , P_{Ar} U(A), 7, 4]
T_2	4	UNDO and CLR[9, T_2 , P_{Cr} U(C), 8, 2]
T_2	2	UNDO and CLR[10, T_2 , -, -, 9, 0]




Recovery (33)

Restart Example 2 (contd.)

- Assumption
 - Crash during Restart

Log File after Crash:



- Analysis Phase
 - As known
- Redo Phase
 - Checking log records of all TA (T_1 , T_2) incl. CLR forward
 - Redo for each CLR



Recovery (34)

Restart Example 2 (contd.)

- Redo Phase (contd.)

TA	Page	Page-LSN	Log-Record-LSN	Action
T_1	P_A	7	3	No REDO
T_2	P_C	4	4	No REDO
T_1	P_B	5	5	No REDO
T_2	P_A	7	7	No REDO
T_2	P_A	7 → 8	8	REDO: U(A)



Recovery (35)

Restart Example 2 (contd.)

- Undo Phase
 - Checking log records of loser TA T_2 backward
 - For each log record Undo is performed and CLR written to log end

TA	Log-Record-LSN	Action
T_2	8	UndoNxtLSN = 4, go to Log Record 4 (Log Record 7 is skipped, since it is already compensated by 8)
T_2	4	UNDO and CLR[9, T_2 , P_C , U(C), 8, 2]
T_2	2	UNDO and CLR[10, T_2 , -, -, 9, 0]



Conclusion (1)

- **Failures**

- Transaction-, System-, Device Failures and Disasters

- **Spectrum of Logging- and Recovery-Mechanisms**

- Entry-Logging outmatches Page-Logging
 - Many DBMS use physiological Logging
 - More flexible recovery within a page
 - Less storage overhead
 - Less I/Os
 - Group Commit



Conclusion (2)

- **Dependencies to other Components**

- Lock granulate must be greater or equal to log granulate
- Atomic
 - Saves DB state of last checkpoint
 - Ensure action consistency
 - Allow logical logging
- Update-in-Place
 - More effective w.r.t. normal operation
 - Low crash probability
 - Require physical logging



Conclusion (3)

- **Basics w.r.t. Update-in-Place**
 - WAL principle: Write Ahead Log for Undo Info
 - Redo Info to be written at the latest at Commit
- **Basics w.r.t. Atomic**
 - WAL principle:
 - TA-related Undo-Info must be written before checkpoint
 - Redo Info to be written at the latest at Commit
- **NoForce**
 - Outmatches force
 - Require checkpoints in order to limit Redo overhead
 - Fuzzy checkpoints cause lowest overhead w.r.t. normal operation



Conclusion (4)

- **Steal**
 - Requires WAL principle
 - Requires Undo actions after crash
- **Restart**
 - Redo action increase page LSNs
 - CLRs for Undo and Rollback actions
 - Restart idempotent

